

Machine Learning and Natural Language Processing Interview Questions

May 3, 2016

- What is POS (Part of Speech) tagging? What is the simplest approach to building a POS tagger that you can imagine?
- How would you train a model that identifies whether the word "Apple" in a sentence belongs to the fruit or the company?
- Design a recommendation system.
- What is latent semantic indexing and where can it be applied?
- How would you build a system that auto corrects text that has been generated by a speech recognition system?
- What is regular grammar? Does this differ in power to a regular expression and if so, in what way?
- How would you build a POS tagger from scratch given a corpus of annotated sentences? How would you deal with unknown words?
- How would you find all the occurrences of quoted text in a news article?
- How would you build a system to translate English text to Greek and vice-versa?
- How would you build a system that automatically groups news articles by subject?
- What are stop words? Describe an application in which stop words should be removed.
- How would you design a model to predict whether a movie review was positive or negative?
- What is entropy? How would you estimate the entropy of the English language?
- What is the TF-IDF score of a word and in what context is this useful?

- How does the PageRank algorithm work?
- What is dependency parsing?
- What are the difficulties in building and using an annotated corpus of text such as the Brown Corpus and what can be done to mitigate them?
- Are you familiar with WordNet or other related linguistic resources?
- Do you have any experience in building ontologies?
- Do you speak any foreign languages?
- What tools for training NLP models (nltk, Apache OpenNLP, GATE, MALLET etc.) have you used?
- Describe approaches to data mining of customer logs.
- What is feature extraction and how would you perform it?
- How would you handle dirty data?
- How do you build a system similar to Netflix?
- How do you implement convolutional neural networks?
- Provide details about LDA and Gibbs Sampling?
- What do you mean by Rejection Sampling?
- Describe a Random Forest.
- Why does one use MSE as a measure of quality. What is the scientific/mathematical reason for the same?
- Given an objective function, calculate the range of its learning rate?
- What is deep learning and what are some of the main characteristics that distinguish it from traditional machine learning?
- What is linear in a generalized linear model?
- Give an example of an application of non-negative matrix factorization
- How would you evaluate the quality of the clusters that are generated by a run of K-means?
- Do you have experience with Spark ML or another platform for building machine learning models using very large datasets?
- What tools and environments have you used to train and assess models?
- What are some good ways for performing feature selection that do not involve exhaustive search?

- On what type of ensemble technique is a random forest based? What particular limitation does it try to address?
- What is a probabilistic graphic model? What is the difference between Markov networks and Bayesian networks?
- What is the EM algorithm? Give a couple of applications
- What methods for dimensionality reduction do you know and how do they compare with each other?
- Explain NLP to a non-technical person.
- What's the use of NLP in machine learning?
- Explain Vector Space Model and its use.
- Explain cosine similarity in a simple way.
- How unstructured text data can be converted into structured data for the purpose of ML models?
- Why and when stop words are removed? In which situation we do not remove them?
- Explain the distances and similarity measures that can be used to compare documents.
- Explain Linear and Logistic Regression and in which scenarios you can use them?
- What is overfitting and what are its disadvantages?